

```
from google.colab import files
uploaded = files.upload()
import pandas as pd
df = pd.read_csv('pd_speech_features.csv')
df.head()
```

Choose Files pd_speech_features.csv

pd_speech_features.csv(text/csv) - 5308926 bytes, last modified: 11/5/2018 - 100% done

Saving pd_speech_features.csv to pd_speech_features (1).csv

	Unnamed: 0	Unnamed: 1	Baseline Features	Unnamed: 3	Unnamed: 4	Unnamed: 5	Unnamed: 6
0	id	gender	PPE	DFA	RPDE	numPulses	numPeriodsPulses
1	0	1	0.85247	0.71826	0.57227	240	239
2	0	1	0.76686	0.69481	0.53966	234	233
3	0	1	0.85083	0.67604	0.58982	232	231
4	1	0	0.41121	0.79672	0.59257	178	177

5 rows × 755 columns

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score
```

```
print(df.shape)
print(df.dtypes)
print(df.describe())
print(df.isnull().sum())
```

```
(757, 755)
Unnamed: 0          object
Unnamed: 1          object
Baseline Features   object
Unnamed: 3          object
Unnamed: 4          object
...
Unnamed: 750         object
Unnamed: 751         object
Unnamed: 752         object
Unnamed: 753         object
Unnamed: 754         object
```

```
Length: 755, dtype: object
      Unnamed: 0 Unnamed: 1 Baseline Features Unnamed: 3 Unnamed: 4 \
count    757        757        757        757        757
unique   253         3         741        746        749
top      0           1       0.82273    0.75192    0.42443
freq     3          390           3           2           2

      Unnamed: 5 Unnamed: 6 Unnamed: 7 Unnamed: 8 Unnamed: 9 ... \
count    757        757        757        757        757 ...
unique   316         320         756        647        359 ...
top      237        236  0.006004477  5.77E-05  0.00076 ...
freq     9           8           2           3           9 ...

      Unnamed: 745 Unnamed: 746 Unnamed: 747 Unnamed: 748 Unnamed: 749 \
count    757        757        757        757        757
unique   750         756         753         754        750
top      1.5382    4.0251    3.0619    3.3603    2.6562
freq     2           2           2           2           2

      Unnamed: 750 Unnamed: 751 Unnamed: 752 Unnamed: 753 Unnamed: 754 \
count    757        757        757        757        757
unique   753         754         754         755         3
top      3.1761    3.1854    4.2391   10.0693      1
freq     2           2           2           2           564

[4 rows x 755 columns]
Unnamed: 0      0
Unnamed: 1      0
Baseline Features  0
Unnamed: 3      0
Unnamed: 4      0
...
Unnamed: 750     0
Unnamed: 751     0
Unnamed: 752     0
Unnamed: 753     0
Unnamed: 754     0
Length: 755, dtype: int64
```

```
df['Unnamed: 1'] = df['Unnamed: 1'].astype('category').cat.codes
X = df.drop(['Unnamed: 0', df.columns[-1]], axis=1)
y = df[df.columns[-1]]
X = X.apply(pd.to_numeric, errors='coerce')
X = X.dropna()
y = y[X.index]
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
```

```
knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train, y_train)
```

▼ KNeighborsClassifier i ?

KNeighborsClassifier(n_neighbors=3)

```
y_pred_knn = knn.predict(X_test)
accuracy_knn = accuracy_score(y_test, y_pred_knn)
print(f'KNN Accuracy: {accuracy_knn:.4f}')
```

KNN Accuracy: 0.9211

```
import pickle
import pandas as pd

# Load the model from the .pkl file
filename = 'knn_optimized_model.pkl'
loaded_model = pickle.load(open(filename, 'rb'))

# Now you can use the loaded_model to make predictions on new data
# For example, if you have new data in a pandas DataFrame called 'new_data_df':
# Make sure 'new_data_df' has the same columns and preprocessing as your training data

# Assuming you have some test data available (like X_test from earlier)
# You would typically use new, unseen data for prediction in a real application
y_pred_loaded = loaded_model.predict(X_test)

# You can compare the predictions with the actual values if you have them
# For example, if you have y_test:
from sklearn.metrics import accuracy_score
accuracy_loaded = accuracy_score(y_test, y_pred_loaded)
print(f'Accuracy of loaded model: {accuracy_loaded:.4f}')

# You can also inspect the loaded model object
print(f'Loaded model type: {type(loaded_model)}')
print(f'Loaded model parameters: {loaded_model.get_params()}')
```

Accuracy of loaded model: 0.9211

Loaded model type: <class 'sklearn.neighbors._classification.KNeighborsClassifier'>
Loaded model parameters: {'algorithm': 'auto', 'leaf_size': 30, 'metric': 'minkowski', 'n_neighbors': 3, 'weights': 'uniform'}

```
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
```

▼ RandomForestClassifier i ?

RandomForestClassifier(random_state=42)

```
import pickle
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import pandas as pd
```

```
# Assuming X_train and y_train are already defined from previous steps

# Train the KNN model with the best n_neighbors (which was found to be 3)
knn_optimized = KNeighborsClassifier(n_neighbors=3)
knn_optimized.fit(X_train, y_train)

# Save the trained KNN model to a .pkl file
filename = 'knn_optimized_model.pkl'
pickle.dump(knn_optimized, open(filename, 'wb'))

print(f"Optimized KNN model saved to {filename}")
```

Optimized KNN model saved to knn_optimized_model.pkl

```
y_pred_rf = rf.predict(X_test)
accuracy_rf = accuracy_score(y_test, y_pred_rf)
print(f'Random Forest Accuracy: {accuracy_rf:.4f}')
```

Random Forest Accuracy: 0.8618

```
svm = SVC(kernel='linear', random_state=42)
svm.fit(X_train, y_train)
```

```
SVC
SVC(kernel='linear', random_state=42)
```

```
y_pred_svm = svm.predict(X_test)
accuracy_svm = accuracy_score(y_test, y_pred_svm)
print(f'SVM Accuracy: {accuracy_svm:.4f}')
```

SVM Accuracy: 0.8618

```
lr = LogisticRegression(random_state=42)
lr.fit(X_train, y_train)
```

```
LogisticRegression
LogisticRegression(random_state=42)
```

```
y_pred_lr = lr.predict(X_test)
accuracy_lr = accuracy_score(y_test, y_pred_lr)
print(f'Logistic Regression Accuracy: {accuracy_lr:.4f}')
```

Logistic Regression Accuracy: 0.8618

```
%pip install xgboost
```

```
Requirement already satisfied: xgboost in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (·)
```

```
Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.12/dist- Requirement already satisfied: scipy in /usr/local/lib/python3.12/dist-packages (
```

```
from xgboost import XGBClassifier
```

```
xgb = XGBClassifier(random_state=42)
y_train_numeric = pd.to_numeric(y_train)
xgb.fit(X_train, y_train_numeric)
```

```
XGBClassifier
```

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              feature_weights=None, gamma=None, grow_policy=None,
              importance_type=None, interaction_constraints=None,
              learning_rate=None, max_bin=None, max_cat_threshold=None,
              max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
              max_leaves=None, min_child_weight=None, missing=np.nan,
              monotone_constraints=None, multi_strategy=None, n_estimators=None,
              n_jobs=None, num_parallel_tree=None, ...)
```

```
y_pred_xgb = xgb.predict(X_test)
accuracy_xgb = accuracy_score(pd.to_numeric(y_test), y_pred_xgb)
print(f'XGBoost Accuracy: {accuracy_xgb:.4f}')
```

```
XGBoost Accuracy: 0.8882
```

```
from sklearn.neural_network import MLPClassifier
mlp = MLPClassifier(hidden_layer_sizes=(128, 64), max_iter=500, random_state=42)
mlp.fit(X_train, y_train)
y_pred_mlp = mlp.predict(X_test)
accuracy_mlp = accuracy_score(y_test, y_pred_mlp)
print(f'MLP Neural Network Accuracy: {accuracy_mlp:.4f}')
```

```
MLP Neural Network Accuracy: 0.8947
```

```
%pip install catboost
from catboost import CatBoostClassifier
cb = CatBoostClassifier(verbose=0, random_state=42)
cb.fit(X_train, y_train)
y_pred_cb = cb.predict(X_test)
accuracy_cb = accuracy_score(y_test, y_pred_cb)
print(f'CatBoost Accuracy: {accuracy_cb:.4f}')
```

```
Collecting catboost
```

```
  Downloading catboost-1.2.8-cp312-cp312-manylinux2014_x86_64.whl.metadata (1.2 kB)
Requirement already satisfied: graphviz in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: numpy<3.0,>=1.16.0 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: pandas>=0.24 in /usr/local/lib/python3.12/dist-packages
```

```
Requirement already satisfied: scipy in /usr/local/lib/python3.12/dist-packages (...)
Requirement already satisfied: plotly in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: six in /usr/local/lib/python3.12/dist-packages (from
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-pac
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-pa
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-pac
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-package
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.12/dist-
Downloading catboost-1.2.8-cp312-cp312-manylinux2014_x86_64.whl (99.2 MB)
99.2/99.2 MB 7.6 MB/s eta 0:00:00
```

Installing collected packages: catboost

Successfully installed catboost-1.2.8

CatBoost Accuracy: 0.8816

```
from lightgbm import LGBMClassifier
lgb = LGBMClassifier(random_state=42)
lgb.fit(X_train, y_train)
y_pred_lgb = lgb.predict(X_test)
accuracy_lgb = accuracy_score(y_test, y_pred_lgb)
print(f'LightGBM Accuracy: {accuracy_lgb:.4f}')
```

```
from sklearn.ensemble import VotingClassifier

voting_clf = VotingClassifier(
    estimators=[
        ('lr', LogisticRegression()),
        ('rf', RandomForestClassifier()),
        ('xgb', XGBClassifier(use_label_encoder=False, eval_metric='logloss'))
    ],
    voting='soft'
)
voting_clf.fit(X_train, y_train)
y_pred_voting = voting_clf.predict(X_test)
accuracy_voting = accuracy_score(y_test, y_pred_voting)
print(f'Voting Ensemble Accuracy: {accuracy_voting:.4f}')

/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: [01
Parameters: { "use_label_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)
Voting Ensemble Accuracy: 0.8882
```

```
from sklearn.ensemble import StackingClassifier

estimators = [
    ('rf', RandomForestClassifier(random_state=42)),
    ('svm', SVC(probability=True, random_state=42)),
    ('xgb', XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_
]

```

```
stacking = StackingClassifier(  
    estimators=estimators,  
    final_estimator=LogisticRegression(),  
    passthrough=True  
)  
stacking.fit(X_train, y_train)  
y_pred_stack = stacking.predict(X_test)  
accuracy_stack = accuracy_score(y_test, y_pred_stack)  
print(f'Stacking Ensemble Accuracy: {accuracy_stack:.4f}')  
  
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: [01  
Parameters: { "use_label_encoder" } are not used.  
  
    bst.update(dtrain, iteration=i, fobj=obj)  
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: [01  
Parameters: { "use_label_encoder" } are not used.  
  
    bst.update(dtrain, iteration=i, fobj=obj)  
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: [01  
Parameters: { "use_label_encoder" } are not used.  
  
    bst.update(dtrain, iteration=i, fobj=obj)  
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: [01  
Parameters: { "use_label_encoder" } are not used.  
  
    bst.update(dtrain, iteration=i, fobj=obj)  
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: [01  
Parameters: { "use_label_encoder" } are not used.  
  
    bst.update(dtrain, iteration=i, fobj=obj)  
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: [01  
Parameters: { "use_label_encoder" } are not used.  
  
    bst.update(dtrain, iteration=i, fobj=obj)  
Stacking Ensemble Accuracy: 0.8750
```

```
from sklearn.ensemble import ExtraTreesClassifier  
et = ExtraTreesClassifier(n_estimators=200, random_state=42)  
et.fit(X_train, y_train)  
y_pred_et = et.predict(X_test)  
accuracy_et = accuracy_score(y_test, y_pred_et)  
print(f'Extra Trees Accuracy: {accuracy_et:.4f}')
```

Extra Trees Accuracy: 0.8684

```
from sklearn.ensemble import GradientBoostingClassifier  
gb = GradientBoostingClassifier(random_state=42)  
gb.fit(X_train, y_train)  
y_pred_gb = gb.predict(X_test)  
accuracy_gb = accuracy_score(y_test, y_pred_gb)  
print(f'Gradient Boosting Accuracy: {accuracy_gb:.4f}')
```

Gradient Boosting Accuracy: 0.8618

```
from sklearn.ensemble import AdaBoostClassifier
ada = AdaBoostClassifier(n_estimators=200, random_state=42)
ada.fit(X_train, y_train)
y_pred_ada = ada.predict(X_test)
accuracy_ada = accuracy_score(y_test, y_pred_ada)
print(f'AdaBoost Accuracy: {accuracy_ada:.4f}')
```

AdaBoost Accuracy: 0.8816

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout

model = Sequential([
    Dense(128, activation='relu', input_shape=(X_train.shape[1],)),
    Dropout(0.3),
    Dense(64, activation='relu'),
    Dense(1, activation='sigmoid')
])

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
model.fit(X_train, pd.to_numeric(y_train), epochs=30, batch_size=32, validation_
```

Epoch 1/30
/usr/local/lib/python3.12/dist-packages/keras/src/layers/core/dense.py:93: UserWarning: super().__init__(activity_regularizer=activity_regularizer, **kwargs)
19/19 ━━━━━━━━ **2s** 21ms/step - accuracy: 0.6541 - loss: 0.6077 - val_

Epoch 2/30
19/19 ━━━━━━━━ **0s** 10ms/step - accuracy: 0.8827 - loss: 0.3133 - val_

Epoch 3/30
19/19 ━━━━━━━━ **0s** 11ms/step - accuracy: 0.9018 - loss: 0.2239 - val_

Epoch 4/30
19/19 ━━━━━━━━ **0s** 20ms/step - accuracy: 0.9379 - loss: 0.1781 - val_

Epoch 5/30
19/19 ━━━━━━━━ **0s** 11ms/step - accuracy: 0.9674 - loss: 0.1232 - val_

Epoch 6/30
19/19 ━━━━━━━━ **0s** 10ms/step - accuracy: 0.9811 - loss: 0.0995 - val_

Epoch 7/30
19/19 ━━━━━━━━ **0s** 9ms/step - accuracy: 0.9892 - loss: 0.0617 - val_a

Epoch 8/30
19/19 ━━━━━━━━ **0s** 9ms/step - accuracy: 0.9957 - loss: 0.0504 - val_a

Epoch 9/30
19/19 ━━━━━━━━ **0s** 10ms/step - accuracy: 0.9905 - loss: 0.0329 - val_

Epoch 10/30
19/19 ━━━━━━━━ **0s** 9ms/step - accuracy: 0.9947 - loss: 0.0278 - val_a

Epoch 11/30
19/19 ━━━━━━━━ **0s** 9ms/step - accuracy: 0.9977 - loss: 0.0225 - val_a

Epoch 12/30
19/19 ━━━━━━━━ **0s** 9ms/step - accuracy: 0.9967 - loss: 0.0213 - val_a

Epoch 13/30
19/19 ━━━━━━━━ **0s** 11ms/step - accuracy: 0.9946 - loss: 0.0307 - val_

Epoch 14/30
19/19 ━━━━━━━━ **0s** 10ms/step - accuracy: 0.9867 - loss: 0.0328 - val_

Epoch 15/30
19/19 ━━━━━━━━ **0s** 9ms/step - accuracy: 0.9984 - loss: 0.0163 - val_a

Epoch 16/30
19/19 ━━━━━━━━ **0s** 12ms/step - accuracy: 0.9961 - loss: 0.0131 - val_

Epoch 17/30

```
19/19 ━━━━━━ 0s 15ms/step - accuracy: 0.9976 - loss: 0.0148 - val_
Epoch 18/30
19/19 ━━━━━━ 1s 14ms/step - accuracy: 1.0000 - loss: 0.0111 - val_
Epoch 19/30
19/19 ━━━━━━ 0s 14ms/step - accuracy: 0.9937 - loss: 0.0279 - val_
Epoch 20/30
19/19 ━━━━━━ 0s 15ms/step - accuracy: 0.9941 - loss: 0.0408 - val_
Epoch 21/30
19/19 ━━━━━━ 0s 14ms/step - accuracy: 0.9988 - loss: 0.0135 - val_
Epoch 22/30
19/19 ━━━━━━ 0s 16ms/step - accuracy: 1.0000 - loss: 0.0043 - val_
Epoch 23/30
19/19 ━━━━━━ 0s 15ms/step - accuracy: 0.9939 - loss: 0.0119 - val_
Epoch 24/30
19/19 ━━━━━━ 0s 9ms/step - accuracy: 1.0000 - loss: 0.0032 - val_a
Epoch 25/30
19/19 ━━━━━━ 0s 9ms/step - accuracy: 1.0000 - loss: 0.0026 - val_a
Epoch 26/30
19/19 ━━━━━━ 0s 10ms/step - accuracy: 1.0000 - loss: 0.0037 - val_
Epoch 27/30
19/19 ━━━━━━ 0s 13ms/step - accuracy: 1.0000 - loss: 0.0021 - val_
Epoch 28/30
```

```
loss, accuracy_nn = model.evaluate(X_test, pd.to_numeric(y_test))
print(f'Neural Network Accuracy: {accuracy_nn:.4f}')
```

```
5/5 ━━━━━━ 0s 8ms/step - accuracy: 0.9107 - loss: 0.3677
Neural Network Accuracy: 0.9079
```

```
%pip install transformers
%pip install hmmlearn
```

```
Requirement already satisfied: transformers in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: huggingface-hub<1.0,>=0.34.0 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: requests in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: tokenizers<=0.23.0,>=0.22.0 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: hf-xet<2.0.0,>=1.1.3 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.12/dist-packages
Collecting hmmlearn
  Downloading hmmlearn-0.3.3-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl
Requirement already satisfied: numpy>=1.10 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: scikit-learn!=0.22.0,>=0.16 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: scipy>=0.19 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages
```

```
  Downloading hmmlearn-0.3.3-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl
Requirement already satisfied: numpy>=1.10 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: scikit-learn!=0.22.0,>=0.16 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: scipy>=0.19 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages
```

166.0/166.0 kB 5.5 MB/s eta 0:00:00

Installing collected packages: hmmlearn
Successfully installed hmmlearn-0.3.3

```
import numpy as np

X_deit = X_scaled.reshape(X_scaled.shape[0], 1, X_scaled.shape[1])

y_deit = pd.to_numeric(y)

print(f'Original X_scaled shape: {X_scaled.shape}')
print(f'Reshaped X_deit shape: {X_deit.shape}')
print(f'y_deit data type: {y_deit.dtype}')

Original X_scaled shape: (756, 753)
Reshaped X_deit shape: (756, 1, 753)
y_deit data type: int64
```

```
from sklearn.model_selection import train_test_split
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Flatten

X_train_deit, X_test_deit, y_train_deit, y_test_deit = train_test_split(
    X_deit, y_deit, test_size=0.2, random_state=42
)

model_deit = Sequential([
    Flatten(input_shape=(X_train_deit.shape[1], X_train_deit.shape[2])),
    Dense(128, activation='relu'),
    Dropout(0.3),
    Dense(64, activation='relu'),
    Dense(1, activation='sigmoid') # binary classification
])

model_deit.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

model_deit.summary()
```

```
/usr/local/lib/python3.12/dist-packages/keras/src/layers/reshaping/flatten.py:37:
    super().__init__(**kwargs)
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 753)	0
dense_7 (Dense)	(None, 128)	96,512
dropout_4 (Dropout)	(None, 128)	0
dense_8 (Dense)	(None, 64)	8,256
dense_9 (Dense)	(None, 1)	65

Total params: 104,833 (409.50 KB)

Trainable params: 104,833 (409.50 KB)

Non-trainable params: 0 (0.00 B)

```
history_deit = model_deit.fit(
    X_train_deit, y_train_deit,
    epochs=30,
    batch_size=32,
    validation_data=(X_test_deit, y_test_deit)
)

19/19 ━━━━━━━━━━ 1s 17ms/step - accuracy: 0.8510 - loss: 0.3342 - val_
Epoch 3/30
19/19 ━━━━━━━━━━ 1s 21ms/step - accuracy: 0.9109 - loss: 0.2191 - val_
Epoch 4/30
19/19 ━━━━━━━━━━ 0s 15ms/step - accuracy: 0.9364 - loss: 0.1857 - val_
Epoch 5/30
19/19 ━━━━━━━━━━ 1s 26ms/step - accuracy: 0.9643 - loss: 0.1278 - val_
Epoch 6/30
19/19 ━━━━━━━━━━ 1s 22ms/step - accuracy: 0.9756 - loss: 0.0980 - val_
Epoch 7/30
19/19 ━━━━━━━━━━ 0s 23ms/step - accuracy: 0.9862 - loss: 0.0612 - val_
Epoch 8/30
19/19 ━━━━━━━━━━ 0s 22ms/step - accuracy: 0.9851 - loss: 0.0623 - val_
Epoch 9/30
19/19 ━━━━━━━━━━ 1s 34ms/step - accuracy: 0.9974 - loss: 0.0302 - val_
Epoch 10/30
19/19 ━━━━━━━━━━ 0s 20ms/step - accuracy: 0.9979 - loss: 0.0356 - val_
Epoch 11/30
19/19 ━━━━━━━━━━ 1s 37ms/step - accuracy: 1.0000 - loss: 0.0193 - val_
Epoch 12/30
19/19 ━━━━━━━━━━ 1s 46ms/step - accuracy: 0.9989 - loss: 0.0155 - val_
Epoch 13/30
19/19 ━━━━━━━━━━ 1s 24ms/step - accuracy: 1.0000 - loss: 0.0123 - val_
Epoch 14/30
19/19 ━━━━━━━━━━ 1s 29ms/step - accuracy: 0.9986 - loss: 0.0118 - val_
Epoch 15/30
19/19 ━━━━━━━━━━ 1s 29ms/step - accuracy: 1.0000 - loss: 0.0130 - val_
```

```
19/19 ━━━━━━ 0s 22ms/step - accuracy: 0.9903 - loss: 0.0166 - val_
Epoch 18/30
19/19 ━━━━━━ 0s 19ms/step - accuracy: 1.0000 - loss: 0.0067 - val_
Epoch 19/30
19/19 ━━━━━━ 1s 24ms/step - accuracy: 1.0000 - loss: 0.0047 - val_
Epoch 20/30
19/19 ━━━━━━ 1s 25ms/step - accuracy: 1.0000 - loss: 0.0034 - val_
Epoch 21/30
19/19 ━━━━━━ 1s 26ms/step - accuracy: 1.0000 - loss: 0.0038 - val_
Epoch 22/30
19/19 ━━━━━━ 1s 35ms/step - accuracy: 1.0000 - loss: 0.0025 - val_
Epoch 23/30
19/19 ━━━━━━ 1s 36ms/step - accuracy: 1.0000 - loss: 0.0020 - val_
Epoch 24/30
19/19 ━━━━━━ 1s 27ms/step - accuracy: 0.9986 - loss: 0.0053 - val_
Epoch 25/30
19/19 ━━━━━━ 1s 28ms/step - accuracy: 1.0000 - loss: 0.0042 - val_
Epoch 26/30
19/19 ━━━━━━ 1s 28ms/step - accuracy: 1.0000 - loss: 0.0037 - val_
Epoch 27/30
19/19 ━━━━━━ 1s 42ms/step - accuracy: 1.0000 - loss: 0.0022 - val_
Epoch 28/30
19/19 ━━━━━━ 0s 22ms/step - accuracy: 0.9993 - loss: 0.0029 - val_
Epoch 29/30
19/19 ━━━━━━ 1s 39ms/step - accuracy: 1.0000 - loss: 0.0040 - val_
Epoch 30/30
19/19 ━━━━━━ 1s 38ms/step - accuracy: 0.9967 - loss: 0.0063 - val_
```

```
loss_deit, accuracy_deit = model_deit.evaluate(X_test_deit, y_test_deit)
print(f'DeiT Model Accuracy: {accuracy_deit:.4f}')
```

```
5/5 ━━━━━━ 0s 64ms/step - accuracy: 0.9041 - loss: 0.4612
DeiT Model Accuracy: 0.8947
```

```
X_astcapsnet = X_scaled.reshape(X_scaled.shape[0], X_scaled.shape[1], 1)

y_astcapsnet = pd.to_numeric(y)

print(f'Original X_scaled shape: {X_scaled.shape}')
print(f'Reshaped X_astcapsnet shape: {X_astcapsnet.shape}')
print(f'y_astcapsnet data type: {y_astcapsnet.dtype}')
```

```
Original X_scaled shape: (756, 753)
Reshaped X_astcapsnet shape: (756, 753, 1)
y_astcapsnet data type: int64
```

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv1D, Flatten, Dense, Dropout, InputLayer
import pandas as pd

X_train_astcapsnet, X_test_astcapsnet, y_train_astcapsnet, y_test_astcapsnet = t
    X_astcapsnet, y_astcapsnet, test_size=0.2, random_state=42
)

model_astcapsnet = Sequential([
    InputLayer(input_shape=(X_train_astcapsnet.shape[1], X_train_astcapsnet.shap
```

```
        Conv1D(filters=64, kernel_size=3, activation='relu'),
        Flatten(),
        Dense(128, activation='relu'),
        Dropout(0.3),
        Dense(64, activation='relu'),
        Dense(1, activation='sigmoid')
    ])

model_astcapsnet.compile(optimizer='adam', loss='binary_crossentropy', metrics=[

history_astcapsnet = model_astcapsnet.fit(
    X_train_astcapsnet, y_train_astcapsnet,
    epochs=30,
    batch_size=32,
    validation_data=(X_test_astcapsnet, y_test_astcapsnet)
)

Epoch 1/30
19/19 ━━━━━━━━━━ 2s 101ms/step - accuracy: 0.8411 - loss: 0.3773 - val
Epoch 3/30
19/19 ━━━━━━━━━━ 3s 101ms/step - accuracy: 0.8973 - loss: 0.2743 - val
Epoch 4/30
19/19 ━━━━━━━━━━ 5s 217ms/step - accuracy: 0.9585 - loss: 0.1365 - val
Epoch 5/30
19/19 ━━━━━━━━━━ 3s 166ms/step - accuracy: 0.9502 - loss: 0.1212 - val
Epoch 6/30
19/19 ━━━━━━━━━━ 3s 167ms/step - accuracy: 0.9870 - loss: 0.0687 - val
Epoch 7/30
19/19 ━━━━━━━━━━ 4s 125ms/step - accuracy: 0.9666 - loss: 0.0592 - val
Epoch 8/30
19/19 ━━━━━━━━━━ 3s 135ms/step - accuracy: 0.9948 - loss: 0.0311 - val
Epoch 9/30
19/19 ━━━━━━━━━━ 2s 100ms/step - accuracy: 0.9949 - loss: 0.0278 - val
Epoch 10/30
19/19 ━━━━━━━━━━ 2s 100ms/step - accuracy: 0.9939 - loss: 0.0301 - val
Epoch 11/30
19/19 ━━━━━━━━━━ 2s 102ms/step - accuracy: 0.9830 - loss: 0.0322 - val
Epoch 12/30
19/19 ━━━━━━━━━━ 2s 101ms/step - accuracy: 0.9972 - loss: 0.0159 - val
Epoch 13/30
19/19 ━━━━━━━━━━ 3s 124ms/step - accuracy: 0.9990 - loss: 0.0100 - val
Epoch 14/30
19/19 ━━━━━━━━━━ 3s 132ms/step - accuracy: 0.9992 - loss: 0.0048 - val
Epoch 15/30
19/19 ━━━━━━━━━━ 2s 106ms/step - accuracy: 1.0000 - loss: 0.0035 - val
Epoch 16/30
19/19 ━━━━━━━━━━ 2s 103ms/step - accuracy: 1.0000 - loss: 0.0035 - val
Epoch 17/30
19/19 ━━━━━━━━━━ 2s 103ms/step - accuracy: 1.0000 - loss: 0.0011 - val
Epoch 18/30
19/19 ━━━━━━━━━━ 2s 102ms/step - accuracy: 1.0000 - loss: 0.0013 - val
Epoch 19/30
19/19 ━━━━━━━━━━ 2s 122ms/step - accuracy: 1.0000 - loss: 0.0014 - val
Epoch 20/30
19/19 ━━━━━━━━━━ 3s 134ms/step - accuracy: 1.0000 - loss: 0.0016 - val
Epoch 21/30
```

```

Epoch 23/30
19/19 2s 101ms/step - accuracy: 0.9948 - loss: 0.0161 - val
Epoch 24/30
19/19 3s 149ms/step - accuracy: 1.0000 - loss: 0.0032 - val
Epoch 25/30
19/19 2s 112ms/step - accuracy: 0.9962 - loss: 0.0110 - val
Epoch 26/30
19/19 2s 103ms/step - accuracy: 0.9938 - loss: 0.0126 - val
Epoch 27/30
19/19 2s 101ms/step - accuracy: 0.9967 - loss: 0.0077 - val
Epoch 28/30
19/19 2s 104ms/step - accuracy: 0.9959 - loss: 0.0074 - val
Epoch 29/30
19/19 2s 101ms/step - accuracy: 1.0000 - loss: 0.0018 - val
Epoch 30/30
19/19 3s 155ms/step - accuracy: 1.0000 - loss: 8.9390e-04 -

```

```

loss_astcapsnet, accuracy_astcapsnet = model_astcapsnet.evaluate(X_test_astcapsnet)
print(f'ASTCapsNet Accuracy: {accuracy_astcapsnet:.4f}')

```

```

5/5 0s 36ms/step - accuracy: 0.9011 - loss: 0.6682
ASTCapsNet Accuracy: 0.8947

```

```

X_hmm = [np.array([seq]) for seq in X_scaled]
y_hmm = pd.to_numeric(y)
print(f'Type of the first sequence in X_hmm: {type(X_hmm[0])}')
print(f'Shape of the first sequence in X_hmm: {X_hmm[0].shape}')
print(f'Data type of y_hmm: {y_hmm.dtype}')

```

```

Type of the first sequence in X_hmm: <class 'numpy.ndarray'>
Shape of the first sequence in X_hmm: (1, 753)
Data type of y_hmm: int64

```

```

from hmmlearn import hmm
X_train_hmm, X_test_hmm, y_train_hmm, y_test_hmm = train_test_split(
    X_hmm, y_hmm, test_size=0.2, random_state=42
)
X_train_concatenated = np.concatenate(X_train_hmm)
lengths_train = [len(x) for x in X_train_hmm]
model_hmm = hmm.GaussianHMM(n_components=2, covariance_type="diag", n_iter=100,
model_hmm.fit(X_train_concatenated, lengths_train)

```

```

WARNING:hmmlearn.base:Some rows of transmat_ have zero sum because no transition -
WARNING:hmmlearn.base:Some rows of transmat_ have zero sum because no transition -
WARNING:hmmlearn.base:Some rows of transmat_ have zero sum because no transition -
WARNING:hmmlearn.base:Some rows of transmat_ have zero sum because no transition -
WARNING:hmmlearn.base:Some rows of transmat_ have zero sum because no transition -
WARNING:hmmlearn.base:Some rows of transmat_ have zero sum because no transition -
WARNING:hmmlearn.base:Some rows of transmat_ have zero sum because no transition -
WARNING:hmmlearn.base:Model is not converging. Current: -546833.9411313558 is not

```

▼
GaussianHMM
 ⓘ

GaussianHMM(n_components=2, n_iter=100, random_state=42)

```
import pandas as pd
from matplotlib import pyplot as plt

model_accuracies = {
    'KNN': accuracy_knn,
    'Random Forest': accuracy_rf,
    'SVM': accuracy_svm,
    'Logistic Regression': accuracy_lr,
    'XGBoost': accuracy_xgb,
    'MLP Neural Network': accuracy_mlp,
    'CatBoost': accuracy_cb,
    'LightGBM': accuracy_lgb,
    'Extra Trees': accuracy_et,
    'Gradient Boosting': accuracy_gb,
    'AdaBoost': accuracy_ada,
    'Voting Ensemble': accuracy_voting,
    'Stacking Ensemble': accuracy_stack,
    'DeiT Model': accuracy_deit,
    'ASTCapsNet': accuracy_astcapsnet,
}

accuracy_df = pd.DataFrame.from_dict(model_accuracies, orient='index', columns=[])
accuracy_df = accuracy_df.sort_values(by='Accuracy', ascending=False)
print("Consolidated Model Accuracy Report:")
display(accuracy_df)
```

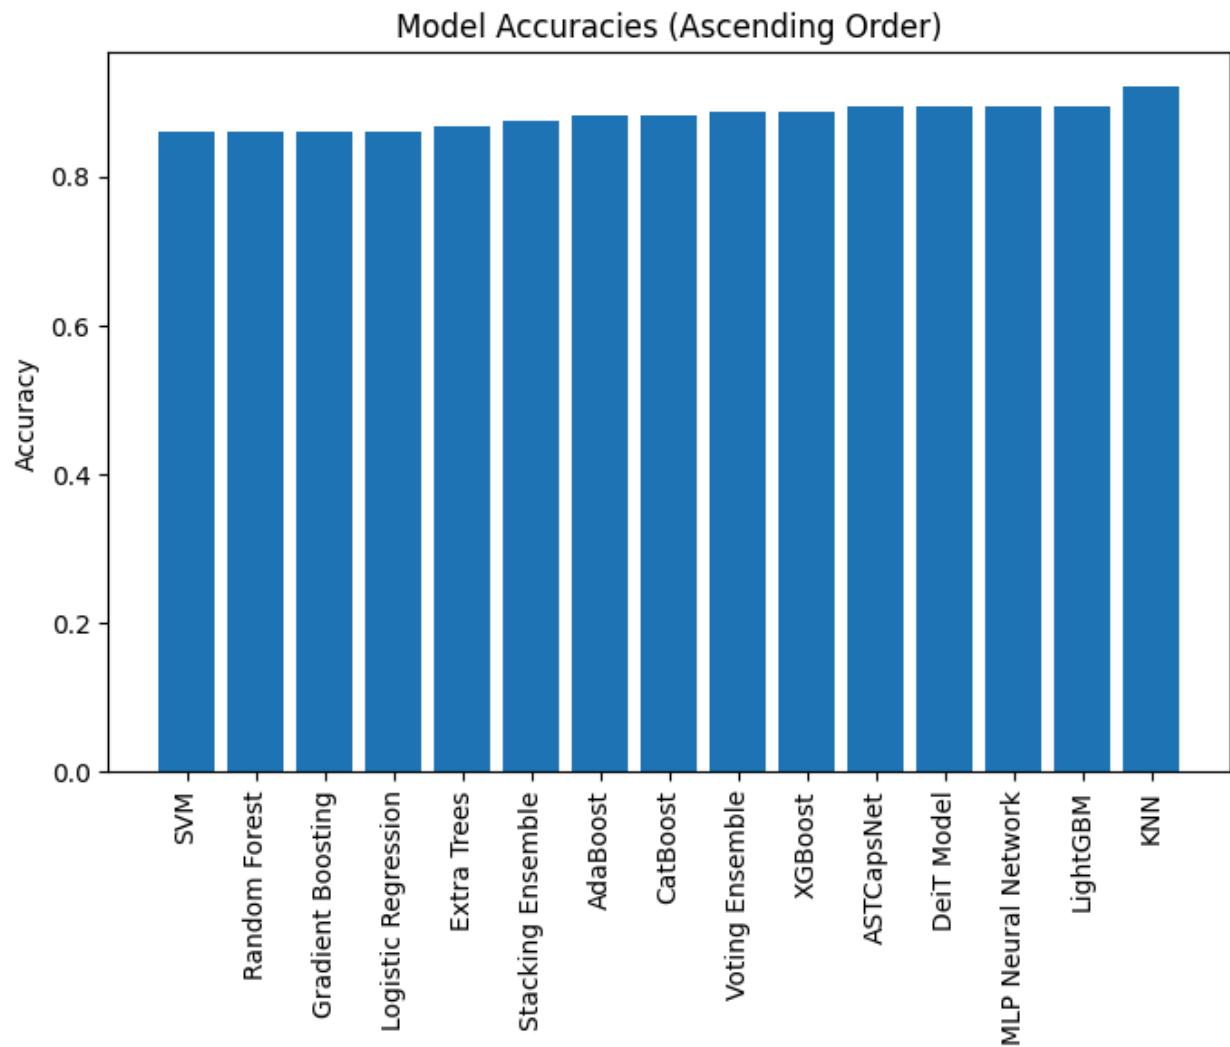
Consolidated Model Accuracy Report:

	Accuracy	
KNN	0.921053	
MLP Neural Network	0.894737	
LightGBM	0.894737	
ASTCapsNet	0.894737	
DeiT Model	0.894737	
Voting Ensemble	0.888158	
XGBoost	0.888158	
AdaBoost	0.881579	
CatBoost	0.881579	
Stacking Ensemble	0.875000	
Extra Trees	0.868421	
Logistic Regression	0.861842	
SVM	0.861842	
Random Forest	0.861842	
Gradient Boosting	0.861842	

Next steps:

[Generate code with accuracy_df](#)[New interactive sheet](#)

```
# Create a bar plot of accuracies in ascending order
accuracy_dfAscending = accuracy_df.sort_values(by='Accuracy', ascending=True)
plt.figure(figsize=(7, 6))
plt.bar(accuracy_dfAscending.index, accuracy_dfAscending['Accuracy'])
plt.xticks(rotation=90)
plt.ylabel('Accuracy')
plt.title('Model Accuracies (Ascending Order)')
plt.tight_layout()
plt.show()
```



Start coding or [generate](#) with AI.